

Week 05 Lectures

PLpgSQL (recap)

1/65

PLpgSQL is a programming language

- containing variables, assignment, conditionals, loops, functions
- combined with database interactions (via SQL)
- functions are stored in the database and invoked from SQL

Example:

```
create or replace function
div(x integer, y integer) returns integer
as $$
declare
    result integer;    -- variable
begin
    if (y <> 0) then    -- conditional
        result := x/y; -- assignment
    else
        result := 0;    -- assignment
    end if;
    return result;
end;
$$ language plpgsql;
```

... PLpgSQL (recap)

2/65

PLpgSQL syntax and control structures

Functions	<pre>create or replace function FunctionName(Params) returns [setof] Type as \$\$ declare Variables begin Code end; \$\$ language plpgsql;</pre>
Assignment	<pre>Variable := Expression select Attrs into Variables ...</pre>
Selection	<pre>if Condition₁ then Statements₁ elsif Condition₂ then Statements₂ ... else Statements_n end if</pre>
Iteration	<pre>loop Statements end loop while Condition loop Statements end loop for IntVariable in Lo..Hi loop ... for RecordVariable in Query loop ...</pre>

Exercise 1: Factorial in PLpgSQL

3/65

Write two PLpgSQL functions ...

- an iterative version of $n!$
- a recursive version of $n!$

Function definition looks like

```
create or replace function
    fac(n integer) returns integer
as $$
...
$$ language plpgsql;
```

Exercise 2: Debugging Output

4/65

Depending on how PostgreSQL is configured

- `raise notice` allows you to display info from a function
- displayed in `psql` window during the function's execution
- usage:
`raise notice StringWith%s, Value1, ... Valuen;`

Write a PLpgSQL function ...

- that takes an integer value n , but returns nothing
- iterates for $1 \dots n$ and prints a message for each value

... PLpgSQL (recap)

5/65

Capturing value(s) from the database, e.g.

```
create or replace function
  full_name(_beer text) returns text
as $$
declare
  _brewer text;
begin
  select manf into _brewer
  from Beers where name = _beer;
  if (not found) then
    return 'No beer called "' || _beer || '"';
  else
    return _brewer || ' ' || _beer;
  end if;
end;
$$ language plpgsql;
```

If query returns multiple tuples, use value(s) in first tuple only.

Exercise 3: Full Names of Beers

6/65

Use the `full_name()` function

- to get names of individual beers
- to handle unknown beers
- to get names of all beers

For the unknown beers case ...

- handle it by returning an 'unknown beer' string
- handle it by raising an `error` exception

Returning Multiple Values

7/65

PLpgSQL functions can return a set of values

- effectively a function returning a table
- values could be atomic \Rightarrow like a single column
- values could be tuples \Rightarrow like a full table

Atomic types, e.g.

`integer`, `float`, `numeric`, `date`, `text`, `varchar(n)`, ...

Tuple types, e.g.

```
create type Point as (x float, y float);
create type Student as (id integer, name text,
                        degree text, wam float, ...);
```

Exercise 4: Functions returning numbers

8/65

Write three functions called `iota()` ...

- each function returns a set of integer values
- `iota(hi)` returns numbers in range 1..*hi*
- `iota(lo,hi)` returns numbers in range lo..*hi*
- `iota(lo,hi,inc)` returns
 *lo, lo+inc, lo+2*inc, lo+3*inc, ..., max ≤ hi*

Functions returning set of *Type* are used in the from clause.

... Returning Multiple Values

9/65

Example function returning a set of tuples

```
create type MyPoint as (x integer, y integer);

create or replace function
  points(n integer, m integer) returns set of MyPoint
as $$
declare
  i integer;  j integer;
  p MyPoint;  -- tuple variable
begin
  for i in 1 .. n loop
    for j in 1 .. m loop
      p.x := i;  p.y := j;
      return next p;
    end loop;
  end loop;
end;
$$ language plpgsql;
```

Query results in PLpgSQL

10/65

Can evaluate a query and iterate through its results

- one tuple at a time, using a `for ... loop`

```
create or replace function
  well_paid(_minsal integer) returns integer
as $$
declare
  nemps integer := 0;
  tuple record;
begin
  for tuple in
    select * from Employees where salary > _minsal
  loop
    nemps := nemps + 1;
  end loop;
  return nemps;
end;
$$ language plpgsql;
```

... Query results in PLpgSQL

11/65

Alternative to the above (but less efficient):

```

create or replace function
    well_paid(_minsal integer) returns integer
as $$
declare
    nemps integer := 0;
    tuple record;
begin
    for tuple in
        select name,salary from Employees
    loop
        if (tuple.salary > _minsal) then
            nemps := nemps + 1;
        end if;
    end loop;
    return nemps;
end;
$$ language plpgsql;

```

INSERT ... RETURNING

12/65

Can capture values from tuples inserted into DB:

```

insert into Table(...) values
(Val1, Val2, ... Valn)
returning ProjectionList into VarList

```

Useful for recording id values generated for serial PKs:

```

declare newid integer; colour text;
...
insert into T(id,a,b,c) values (default,2,3,'red')
returning id,c into newid,colour;
-- id contains the primary key value
-- for the new tuple T(?,2,3,'red')

```

Exceptions

13/65

PLpgSQL supports exception handling via

```

begin
    Statements...
exception
    when Exceptions1 then
        StatementsForHandler1
    when Exceptions2 then
        StatementsForHandler2
    ...
end;

```

Each *Exceptions_i* is an OR list of exception names, e.g.

division_by_zero OR floating_point_exception OR ...

A list of exceptions is in Appendix A of the PostgreSQL Manual.

... Exceptions

14/65

When an exception occurs:

- control is transferred to the relevant exception handling code
- all database changes so far in this transaction are undone
- all function variables retain their current values
- handler executes and then transaction aborts (and function exits)

If no handler in current scope, exception passed to next outer level.

Default exception handlers at outermost level exit and log error.

... Exceptions

15/65

Example of exception handling:

```
-- table T contains one tuple ('Tom','Jones')
declare
  x integer := 3;
begin
  update T set firstname = 'Joe'
  where lastname = 'Jones';
  -- table T now contains ('Joe','Jones')
  x := x + 1;
  y := x / 0;
exception
  when division_by_zero then
    -- update on T is rolled back to ('Tom','Jones')
    raise notice 'caught division_by_zero';
    return x; -- return may or may not work here
    -- if it does, value returned is 4
end;
```

... Exceptions

16/65

The raise operator can generate server log entries, e.g.

```
raise debug 'Simple message';
raise notice 'User = %',user_id;
raise exception 'Fatal: value was %',value;
```

There are several levels of severity:

- DEBUG, LOG, INFO, NOTICE, WARNING, and EXCEPTION
- not all severities generate a message to the client

Your CSE server log is the file `/srvr/YOU/pgsql/Log`

Server logs can grow *very* large; delete when you shut your server down

Dynamically Generated Queries

17/65

EXECUTE takes a string and executes it as an SQL query.

Examples:

```
execute 'select * from Employees';
execute 'select * from '||'Employees';
execute 'select * from '||quote_ident($1);
execute 'delete from Accounts '||
  'where holder='||quote_literal($1);
```

Can be used in any context where an SQL query is expected

This mechanism allows us to *construct* queries "on the fly".

... Dynamically Generated Queries

18/65

Example: a wrapper for updating a single text field

```
create or replace function
  set(_tab text, _attr text, _val text) returns void
as $$
declare
  query text;
begin
```

```

query := 'update ' || quote_ident(_tab);
query := query || ' SET ' || quote_ident(_attr);
query := query || ' = ' || quote_literal(_val);
EXECUTE query;
end; $$ language plpgsql;

```

which could be used as e.g.

```
select set('branches','address','Beach St.');
```

... Dynamically Generated Queries

19/65

One limitation of EXECUTE:

- cannot use `select into` inside dynamic queries

Needs to be expressed instead as:

```

declare tuple R%rowtype; n int;
execute 'select * from R where id=' || n into tuple;
-- or
declare x int; y int; z text;
execute 'select a,b,c from R where id=' || n into x,y,z;

```

Notes:

- if query returns multiple tuples, first one is stored
- if query returns zero tuples, all nulls are stored

Functions vs Views

20/65

A difference between views and functions returning a SETOF:

- CREATE VIEW produces a "virtual" table definition
- SETOF functions require an existing tuple type

In examples above, we used existing Employees tuple type.

In general, you need to define the tuple return type via

```
create type TupleType as ( attr1 type1, ... attrn typen );
```

Other major differences between setof functions and views ...

- functions have parameters; views don't (although where might help)
- functions are "run-time" objects; views are interpolated into queries

... Functions vs Views

21/65

Example of function returning setof tuples ...

```

create type EmpInfo as (name text, pay integer);

create or replace function
  richEmps(_minsal integer) returns setof EmpInfo
as $$
declare
  emp record;   info EmpInfo;
begin
  for emp in
    select * from Employees where salary > _minsal
  loop
    info.name := emp.name;
    info.pay := emp.salary;
    return next info;
  end loop;
end; $$ language plpgsql;

```

Using the function ...

```
select * from richEmps(100000);
```

versus a view

```
create view richEmps(name,pay) as
select name, salary from Employees where salary > 100000;

select * from richEmps; -- but no scope for different salary
```

versus an SQL function

```
create function
    richEmps(_minsal integer) returns setof EmpInfo
as $$
select name, salary from Employees where salary > _minsal;
$$ language sql;
```

Aggregates

Aggregates reduce a collection of values into a single result.

Examples: `count(Tuples)`, `sum(Numbers)`, `max(AnyOrderedType)`

The action of an aggregate function can be viewed as:

```
State = initial state
for each item V {
    # update State to include V
    State = updateState(State, V)
}
return makeFinal(State)
```

... Aggregates

Aggregates are commonly used with `GROUP BY`.

In that context, they "summarise" each group.

Example:

<pre>R a b c ---+---+--- 1 2 x 1 3 y 2 2 z 2 1 a 2 3 b</pre>	<pre>select a,sum(b),count(*) from R group by a a sum count ---+---+--- 1 5 2 2 6 3</pre>
--	--

User-defined Aggregates

SQL standard does not specify user-defined aggregates.

But PostgreSQL provides a mechanism for defining them.

To define a new aggregate, first need to supply:

- *BaseType* ... type of input values
- *StateType* ... type of intermediate states
- state mapping function: *sfunc(state,value) → newState*
- [optionally] an initial state value (defaults to null)
- [optionally] final function: *ffunc(state) → result*

... User-defined Aggregates

27/65

New aggregates defined using CREATE AGGREGATE statement:

```
CREATE AGGREGATE AggName(BaseType) (
    sfunc      = UpdateStateFunction,
    stype      = StateType,
    initcond   = InitialValue,
    finalfunc  = MakeFinalFunction,
    sortop     = OrderingOperator
);
```

- *initcond* (type *StateType*) is optional; defaults to NULL
- *finalfunc* is optional; defaults to identity function
- *sortop* is optional; needed for min/max-type aggregates

... User-defined Aggregates

28/65

Example: defining the count aggregate (roughly)

```
create aggregate myCount(anyelement) (
    stype      = int,      -- the accumulator type
    initcond   = 0,        -- initial accumulator value
    sfunc      = oneMore   -- increment function
);

create function
    oneMore(sum int, x anyelement) returns int
as $$
begin return sum + 1; end;
$$ language plpgsql;
```

... User-defined Aggregates

29/65

Example: sum2 sums two columns of integers

```
create type IntPair as (x int, y int);

create function
    AddPair(sum int, p IntPair) returns int
as $$
begin return p.x + p.y + sum; end;
$$ language plpgsql;

create aggregate sum2(IntPair) (
    stype      = int,
    initcond   = 0,
    sfunc      = AddPair
);
```

Exercise 5: Product Aggregate

30/65

PostgreSQL has many aggregates (e.g. sum, count, ...)

But it doesn't have a product aggregate.

Implement a prod aggregate that

- computes the product of values in a column of integer data

Usage:

```
select prod(*) from iota(5);
prod
-----
120
```

But we need: `select prod(iota::integer) from iota(5)`

Exercise 6: String Concatenation Aggregate

31/65

Define a `concat` aggregate that

- takes a column of string values
- returns a comma-separated string of values

Example:

```
select count(*), concat(name) from Employee;
-- returns e.g.
count |          concat
-----+-----
4 | John,Jane,David,Phil
```

Use it to get a list of beers liked by each drinker.

Constraints

Constraints

33/65

So far, we have considered several kinds of constraints:

- **attribute** (column) constraints
- **relation** (table) constraints
- **referential integrity** constraints

Examples:

```
create table Employee (
  id      integer primary key,
  name    varchar(40),
  salary  real,
  age     integer check (age > 15),
  worksIn integer
           references Department(id),
  constraint PayOk check (salary > age*1000)
);
```

... Constraints

34/65

Column and table constraints ensure validity of one table.

Ref. integrity constraints ensure connections between tables are valid.

However, specifying validity of entire database often requires constraints involving multiple tables.

Simple example (from banking domain):

```
for all Branches b
  b.assets == (select sum(acct.balance)
               from   Accounts acct
               where  acct.branch = b.location)
```

i.e. assets of a branch is sum of balances of accounts held at that branch

Assertions are schema-level constraints

- typically involving multiple tables
- expressing a condition that must hold at all times
- need to be checked on each change to relevant tables
- if change would cause check to fail, reject change

SQL syntax for assertions:

```
CREATE ASSERTION name CHECK (condition)
```

The *condition* is expressed as "there are no violations in the database"

Implementation: ask a query to find all the violations; check for empty result

... Assertions

36/65

Example: #students in any UNSW course must be < 10000

```
create assertion ClassSizeConstraint check (  
  not exists (  
    select c.id from Courses c, CourseEnrolments e  
    where c.id = e.course  
    group by c.id having count(e.student) > 9999  
  )  
);
```

Needs to be checked after *every* change to either Course or CourseEnrolment

... Assertions

37/65

Example: assets of branch = sum of its account balances

```
create assertion AssetsCheck check (  
  not exists (  
    select branchName from Branches b  
    where b.assets <>  
          (select sum(a.balance) from Accounts a  
           where a.branch = b.location)  
    )  
);
```

Needs to be checked after *every* change to either Branch or Account

... Assertions

38/65

On each update, it is expensive

- to determine which assertions need to be checked
- to run the queries which check the assertions

A database with many assertions would be way too slow.

So, most RDBMSs do not implement general assertions.

Typically, *triggers* are provided as

- a lightweight mechanism for dealing with assertions
 - a general event-based programming tool for databases
-

Triggers

Triggers

Triggers are

- procedures stored in the database
- activated in response to database events (e.g. updates)

Examples of uses for triggers:

- maintaining summary data
- checking schema-level constraints (assertions) on update
- performing multi-table updates (to maintain assertions)

... Triggers

41/65

Triggers provide event-condition-action (ECA) programming:

- an *event* activates the trigger
- on activation, the trigger checks a *condition*
- if the condition holds, a procedure is executed (the *action*)

Some typical variations on this:

- execute the action before, after or instead of the triggering event
- can refer to both old and new values of updated tuples
- can limit updates to a particular set of attributes
- perform action: for each modified tuple, once for all modified tuples

... Triggers

42/65

SQL "standard" syntax for defining triggers:

```
CREATE TRIGGER TriggerName
{AFTER|BEFORE} Event1 [ OR Event2 ... ]
[ FOR EACH ROW ]
ON TableName
[ WHEN ( Condition ) ]
Block of Procedural/SQL Code ;
```

Possible *Events* are INSERT, DELETE, UPDATE.

FOR EACH ROW clause ...

- if present, code is executed on each modified tuple
- if not present, code is executed once after all tuples are modified, just before changes are finally COMMITED

Trigger Semantics

43/65

Triggers can be activated BEFORE or AFTER the event.

If activated BEFORE, can affect the change that occurs:

- NEW contains "proposed" value of changed tuple
- modifying NEW causes a different value to be placed in DB

If activated AFTER, the effects of the event are visible:

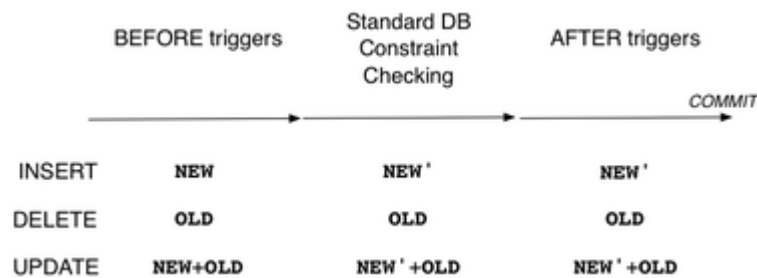
- NEW contains the current value of the changed tuple
- OLD contains the previous value of the changed tuple
- constraint-checking has been done for NEW

Note: OLD does not exist for insertion; NEW does not exist for deletion.

... Trigger Semantics

44/65

Sequence of activities during database update:



Reminder: BEFORE trigger can modify value of new tuple

... Trigger Semantics

45/65

Consider two triggers and an INSERT statement

```
create trigger X before insert on T Code1;
create trigger Y after insert on T Code2;
insert into T values (a,b,c,...);
```

Sequence of events:

- execute Code1 for trigger X
- code has access to (a,b,c,...) via NEW
- code typically checks the values of a,b,c,...
- code can modify values of a,b,c,... in NEW
- DBMS does constraint checking as if NEW is inserted
- if fails any checking, abort insertion and rollback
- execute Code2 for trigger Y
- code has access to final version of tuple via NEW
- code typically does final checking, or modifies other tables in database to ensure constraints are satisfied

Reminder: there is no OLD tuple for an INSERT trigger.

... Trigger Semantics

46/65

Consider two triggers and an UPDATE statement

```
create trigger X before update on T Code1;
create trigger Y after update on T Code2;
update T set b=j,c=k where a=m;
```

Sequence of events:

- execute Code1 for trigger X
- code has access to current version of tuple via OLD
- code has access to updated version of tuple via NEW
- code typically checks new values of b,c,...
- code can modify values of a,b,c,... in NEW
- do constraint checking as if NEW has replaced OLD
- if fails any checking, abort update and rollback
- execute Code2 for trigger Y
- code has access to final version of tuple via NEW
- code typically does final checking, or modifies other tables in database to ensure constraints are satisfied

Reminder: both OLD and NEW exist in UPDATE triggers.

... Trigger Semantics

47/65

Consider two triggers and an DELETE statement

```
create trigger X before delete on T Code1;
create trigger Y after delete on T Code2;
delete from T where a=m;
```

Sequence of events:

- execute Code1 for trigger X
- code has access to (a,b,c,...) via OLD
- code typically checks the values of a,b,c,...
- DBMS does constraint checking as if OLD is removed
- if fails any checking, abort deletion (restore OLD)
- execute Code2 for trigger Y
- code has access to about-to-be-deleted tuple via OLD
- code typically does final checking, or modifies other tables in database to ensure constraints are satisfied

Reminder: tuple NEW does not exist in DELETE triggers.

Triggers in PostgreSQL

48/65

PostgreSQL triggers provide a mechanism for

- INSERT, DELETE or UPDATE events
- to automatically activate PLpgSQL functions

Syntax for PostgreSQL trigger definition:

```
CREATE TRIGGER TriggerName
{AFTER|BEFORE} Event1 [OR Event2 ...]
ON TableName
[ WHEN ( Condition ) ]
FOR EACH {ROW|STATEMENT}
EXECUTE PROCEDURE FunctionName(args...);
```

... Triggers in PostgreSQL

49/65

There is no restriction on what code can go in the function.

However a BEFORE function must contain one of:

```
RETURN old;      or      RETURN new;
```

depending on which version of the tuple is to be used.

If BEFORE trigger returns old, no change occurs.

If exception is raised in trigger function, no change occurs.

Trigger Example #1

50/65

Consider a database of people in the USA:

```
create table Person (
  id      integer primary key,
  ssn     varchar(11) unique,
  ... e.g. family, given, street, town ...
  state   char(2), ...
);
create table States (
  id      integer primary key,
  code    char(2) unique,
  ... e.g. name, area, population, flag ...
);
```

Constraint: `Person.state ∈ (select code from States)`, or
`exists (select id from States where code=Person.state)`

... Trigger Example #1

51/65

Example: ensure that only valid state codes are used:

```
create trigger checkState before insert or update
```

```

on Person for each row execute procedure checkState();

create function checkState() returns trigger as $$
begin
    -- normalise the user-supplied value
    new.state = upper(trim(new.state));
    if (new.state !~ '^[A-Z][A-Z]$') then
        raise exception 'Code must be two alpha chars';
    end if;
    -- implement referential integrity check
    select * from States where code=new.state;
    if (not found) then
        raise exception 'Invalid code %',new.state;
    end if;
    return new;
end;
$$ language plpgsql;

```

... Trigger Example #1

52/65

Examples of how this trigger would behave:

```

insert into Person
    values('John',..., 'Calif.',...);
-- fails with 'Statecode must be two alpha chars'

insert into Person
    values('Jane',..., 'NY',...);
-- insert succeeds; Jane lives in New York

update Person
    set town='Sunnyvale', state='CA'
    where name='Dave';
-- update succeeds; Dave moves to California

update Person
    set state='OZ' where name='Pete';
-- fails with 'Invalid state code OZ'

```

Trigger Example #2

53/65

Example: department salary totals

Scenario:

```

Employee(id, name, address, dept, salary, ...)
Department(id, name, manager, totSal, ...)

```

An assertion that we wish to maintain:

```

create assertion TotalSalary check (
    not exists (
        select d.id from Department d
        where d.totSal <>
            (select sum(e.salary)
             from Employee e
             where e.dept = d.id)
    )
)

```

... Trigger Example #2

54/65

Events that might affect the validity of the database

- a new employee starts work in some department
- an employee gets a rise in salary
- an employee changes from one department to another
- an employee leaves the company

A single assertion could check for this after each change.

With triggers, we have to program each case separately.

Each program implements updates to *ensure* assertion holds.

... Trigger Example #2

55/65

Implement the Employee update triggers from above in PostgreSQL:

Case 1: new employees arrive

```
create trigger TotalSalary1
after insert on Employees
for each row execute procedure totalSalary1();

create function totalSalary1() returns trigger
as $$
begin
    if (new.dept is not null) then
        update Department
        set     totSal = totSal + new.salary
        where  Department.id = new.dept;
    end if;
    return new;
end;
$$ language plpgsql;
```

... Trigger Example #2

56/65

Case 2: employees change departments/salaries

```
create trigger TotalSalary2
after update on Employee
for each row execute procedure totalSalary2();

create function totalSalary2() returns trigger
as $$
begin
    update Department
    set     totSal = totSal + new.salary
    where  Department.id = new.dept;
    update Department
    set     totSal = totSal - old.salary
    where  Department.id = old.dept;
    return new;
end;
$$ language plpgsql;
```

... Trigger Example #2

57/65

Case 3: employees leave

```
create trigger TotalSalary3
after delete on Employee
for each row execute procedure totalSalary3();

create function totalSalary3() returns trigger
as $$
begin
    if (old.dept is not null) then
        update Department
        set     totSal = totSal - old.salary
        where  Department.id = old.dept;
    end if;
    return old;
end;
```

```
end;
$$ language plpgsql;
```

Exercise 7: Triggers (1)

58/65

Requirement: maintain assets in bank branches

- each branch has assets based on the accounts held there
- whenever an account changes, the assets of the corresponding branch should be updated to reflect this change

Some possible changes:

- a new account is opened
- the amount of money in an account changes
- an account moves from one branch to another
- an account is closed

Implement triggers to maintain `Branch.assets`

Exercise 8: Triggers (2)

59/65

Consider a simple airline flights/bookings database:

```
Airports(id, code, name, city)
Planes(id, craft, nseats)
Flights(id, fltNum, plane, source, dest
        departs, arrives, price, seatsAvail)
Passengers(id, name, address, phone)
Bookings(pax, flight, paid)
```

Write triggers to ensure that `Flights.seatsAvail` is consistent with number of `Bookings` on that flight.

Assume that we never `UPDATE` a booking (only insert/delete)

Catalogs

DBMS Catalog

61/65

DBMSs store ...

- data (tuples, organised into tables)
- stored procedures (e.g. PLpgSQL functions)
- indexes (to provide efficient access to data)
- meta-data (information giving the structure of the data)

The latter is stored in the *system catalog*

A standard `information_schema` exists for describing meta-data

But was developed long after DBMSs had implemented their own catalogs

PostgreSQL has both: PG catalog Ch.52, Information Schema Ch.37

PostgreSQL Catalog

62/65

Catalog = meta-data = tables describing DB objects

PostgreSQL catalog is accessible via **pg_XXX** tables/views:

```
pg_roles(oid, rolname, rolsuper, ...)
```

```
pg_namespace(oid, nspname, nspowner, nspacl)
```



```
pg_database(oid, datname, datdba, ..., datacl)
pg_class(oid, relname, relnamespace, reltype, ...)
pg_attribute(oid, attrrelid, attrname, atttypid, ...)
pg_type(oid, typname, typnamespace, typowner, ...)
```

Catalog tables use `oid` field for PKs/FKs

Standard-format catalog also available, via `information_schema`

Exercise 9: List of Databases

63/65

Use the `pg_catalog` to implement simplified `psql -l`

```
$ psql -l | cut -c1-23
```

Name	Owner
a1	jas
a2	jas
acad	jas
accreditation	jas
aims	jas
assess	jas
bank	jas
beer	jas
cmap	jas
company	jas
cse	jas
...	

Exercise 10: Size of Database Tables

64/65

Implement a `dbPop()` function that ...

- lists all of the tables in the public schema
- counts the number of tuples in each table

The function is defined as

```
function dbPop() returns setof PopulationRecord
```

where each return tuple has the type

```
PopulationRecord(table_name, n_records)
```

Exercise 11: Simple Schema Dump

65/65

Implement a function `dbSchema()` that gives a list of tables in the public schema and their attributes.

The function is defined as

```
function dbSchema()
    returns setof SchemaRecord
```

where each return tuple has the type

```
SchemaRecord(table_name, attributes)
```

Attributes = comma-separated list of names, in order.
